

Per-step GD loss at convergence — reduced3

